

Tutorial 09 — Graph Traversal

CS2040C Semester 2 2018/2019

By Jin Zhe, adapted from slides by Ranald, AY1819 S2 Tutor

Notations

For this tutorial, unless specified otherwise

- G : The entire graph
- V : The total number of vertices in G
- E : The total number of edges in G
- u, v : vertices in G
- e : an edge
- CC: A connected component in G

Review Harder Topics

Question 1

In increasing order of difficulty, the following topics are usually found to be challenging for students:

1. <https://visualgo.net/en/dfsbfbs?slide=7-1> to 7-3: Detecting back edge/cycle in the graph
2. <https://visualgo.net/en/dfsbfbs?slide=7-10> to 7-11: Toposort, revisited in Q3 & Q4
3. <https://visualgo.net/en/dfsbfbs?slide=7-6> to 7-9: Check your understanding about the $O(V(V+E))$ versus just $O(V+E)$ analysis again
4. <https://visualgo.net/en/dfsbfbs?slide=8>: All other more advanced graph traversal topics not under the main focus of CS2040/C are optional and will not be part of CS2040/C final assessment

DFS Complexity Analysis — Test yourself!

Claim: The time complexity of DFS is $O(V+E)$.

True or False?

DFS Complexity Analysis — Test yourself!

Claim: The time complexity of DFS is $O(V+E)$.

True!

DFS Complexity Analysis—Why $O(V+E)$?

Recall that time complexity of DFS on binary trees with V vertices is $3 \times V = O(V)$?

3 for each vertex because:

- 1 visitation operation
- 2 comparisons on its 2 children

DFS Complexity Analysis—Why $O(V+E)$?

It's easy to see why this is the case when we write binary tree DFS like this:

```
void dfs(vertex v) {  
    visited[v.id] = true;           // 1 visit operation  
    /* If have left child and it's unvisited */  
    if (v.left && !visited[v.id])   // Treat as 1 comparison  
        dfs(v.left);  
    /* If have right child and it's unvisited */  
    if (v.right && !visited[v.right]) // Treat as 1 comparison  
        dfs(v.right);  
}
```


DFS Complexity Analysis—Why $O(V+E)$?

So for general graph DFS, a vertex v entails 1 visit operation and e_v comparisons where e_v is its number of outgoing edges.

```
void dfs(int v_id) {  
    visited[v_id] = true;           // 1 visit operation  
    for (auto &nb_id : AL[v_id]) { //  $e_{v\_id}$  neighbours  
        if (!visited[nb_id])      // 1 comparison  
            dfs(nb_id);  
    }  
}
```

DFS Complexity Analysis—Why $O(V+E)$?

So in general, running DFS on a connected graph with V vertices and E edges incurs time complexity:

$$\begin{aligned} & 1 + e_1 + 1 + e_2 + 1 + e_3 + \dots + 1 + e_V \\ = & \underbrace{1 + 1 + \dots + 1}_V + \underbrace{e_1 + e_2 + \dots + e_V}_E && \text{Grouping 1's and } E\text{'s separately} \\ = & V + E \\ = & O(V + E) \quad Q.E.D \end{aligned}$$

DFS Complexity Analysis — Test yourself!

Claim: When finding connected components (CC), we perform DFS once per connected component.

True or False?

DFS Complexity Analysis — Test yourself!

Claim: When finding connected components (CC), we perform DFS once per connected component.

True!

We will see why this is the case when we see the flood-fill examples.

DFS Complexity Analysis — Test yourself!

Claim: There are at most V CCs in a graph

True or False?

DFS Complexity Analysis — Test yourself!

Claim: There are at most V CCs in a graph

True!

Extreme case when every vertex is an isolated component.
i.e. graph has no edges

DFS Complexity Analysis — Test yourself!

Claim: When we run DFS to detect every CC, the total time complexity is $V \times O(V+E) = O(V^2 + VE)$ because we run DFS on every vertex in the graph.

True or False?


DFS Complexity Analysis — Test yourself!

Claim: When we run DFS to detect every CC, the total time complexity is $V \times O(V+E) = O(V^2 + VE)$ because we run DFS on every vertex in the graph.


False!

It is still $O(V+E)$! Why? Recall that we will only run DFS once per CC. After DFS ran in a CC, all its vertices will be marked as visited.


DFS Complexity Analysis—Connected Components



CC 1:
 V_1 vertices, E_1 edges



CC 2:
 V_2 vertices, E_2 edges



CC 3:
 V_3 vertices, E_3 edges

Total complexity of running
DFS to detect every CC:

DFS on CC 1: $O(V_1 + E_1) +$

DFS on CC 2: $O(V_2 + E_2) +$

DFS on CC 3: $O(V_3 + E_3) =$

$O(V + E)$

Relevant discussion

I have V integers distributed across C arrays.

$$1 \leq C \leq V$$

Sort each array using $O(N \log N)$ sort

What is the total time complexity?

$$O(V \log V)$$


BFS Complexity Analysis—Test yourself!

Does all the previous analysis for DFS also applies to BFS?

DFS Spanning Tree

[Credits: SG IOI Training Team]

DFS visitation order entails an implicit **directed spanning tree** rooted at the source vertex which DFS ran from.




DFS Spanning Tree

[Credits: SG IOI Training Team]

For instance, we can have the right DFS spanning tree if we run DFS starting from the topmost vertex in the graph.

Note: Dashed lines represent edges untraversed by DFS. i.e edges not in the spanning tree




DFS Spanning Tree

[Credits: SG IOI Training Team]

The untraversed edges by DFS are classified into:

- **Forward edges**
- **Cross edges**
- **Back edges**

Note: You don't have to know the significance of **Forward** and **Cross** edges in CS2040C.



DFS Application — Cycle Detection

What you need to know in CS2040C:

- The only edges that can create cycles are **back edges**
- **Back edges** always point from a vertex to one of its ancestors
- Therefore to detect cycles in graph
 - Run DFS
 - Check for **back edges**
 - If exists, then there is at least a cycle
 - If not exists, then there are no cycles

DFS Application — Vertex States

During DFS, we maintain 3 visitation states for each vertex:

1. **Unvisited** : Vertex is still “untouched” by DFS
2. **Visiting** : Vertex is “touched” but not yet completed by DFS
3. **Visited** : Vertex is “touched” and completed by DFS

At any level of the DFS recursion:

- Ancestors of current vertex will be in **Visiting** state
- “Touched” non-ancestors will be in **Visited** state

DFS Application — Back Edge Detection

Back edge always point from current vertex in DFS recursion to one of its ancestors.


Therefore to detect **back edge**, it suffices to check if the edge go from **Visiting** state → **Visiting** state.

DFS Application — Cycle Detection

Why do we need **Visiting** state?

In other words why can't **back edges** go from **Visited** → **Visited**?

To find out, we shall trace out DFS on the graph starting from vertex 2.



Click [here](#) for the graph in VisuAlgo


DFS Application — Cycle Detection

Without **Visiting** state:

Vertex 5: **Current Vertex**

Vertex 2: **Visited**

Since 2 is **Visited** state, we conclude that $5 \rightarrow 2$ is a **back edge**, which is indeed the case. So far so good...



DFS Application — Cycle Detection


Without **Visiting** state:

Vertex 7: **Current Vertex**

Vertex 4: **Visited**

Since 4 is **Visited** state, we conclude that $7 \rightarrow 4$ is a **back edge**.

This is wrong! It's a **cross edge**!




DFS Application — Cycle Detection

Without **Visiting** state:

Thus it is evident that we cannot differentiate between **back edges** and **cross edges** with just 2 states.

Next we shall see what happens if we introduce the **Visiting** state!




DFS Application—Cycle Detection

With **Visiting** state:

Vertex 5: **Current Vertex**

Vertex 2: **Visiting**

Since 2 is **Visiting** state, we conclude that $5 \rightarrow 2$ is a **back edge**, which is indeed the case!




DFS Application — Cycle Detection

With **Visiting** state:

Vertex 7: **Current Vertex**

Vertex 4: **Visited**


Since 4 is **Visited** state, we conclude that $7 \rightarrow 4$ is a **cross edge**, which is indeed the case! We did not make the same mistake as before.



DFS Application — Cycle Detection

With **Visiting** state:

Thus we need both **Visiting** and **Visited** states in order to differentiate between **back edges** and **cross edges**!



DFS Application—Cycle Detection

The following is one way to implement the modified DFS subroutine for cycle detection in a graph:

```
void detect_cycle(int v_id) {  
    if (is_visited(nb_id)) return;  
    mark_visiting(v_id);           // Mark visiting  
    for (auto &nb_id : AL[v_id]) {  
        if (is_visiting(nb_id)) // If backedge  
            has_cycle = true;  
        else  
            detect_cycle(nb_id);  
    }  
    mark_visited(v_id);           // Mark visited  
}
```

DFS Application—Cycle Detection

Our main procedure might look something like this:

```
...  
bool has_cycle = false;  
/* Iterate over each vertex v_id in graph */  
for (int v_id = 0; v_id < AL.size(); ++v_id) {  
    detect_cycle(v_id);    // Detect on every vertex v_id  
}  
/* Do something about the detection */  
if (has_cycle)  
    cout << "This graph has cycle(s)!" << endl;  
...
```

DFS/BFS Applications

Question 2

You have already seen several examples of DFS/BFS applications from the questions discussed in Week 10

1. <https://nus.kattis.com/problems/countingstars>
 - Easy flood-filling/finding CCs
2. <https://nus.kattis.com/problems/reachableroads>
 - Easy DFS/BFS application; finding CCs too
3. <https://nus.kattis.com/problems/runningmom>
 - Back edge/cycle detection problem

Example Problem 1: Counting Stars

You are given a m by n grid of characters representing the night sky.

Each cell is either

- #: black pixels
- -: white pixels

White pixels that are adjacent vertically or horizontally are part of the same star.

You are to count the number of stars.

Counting Stars—Approach

Whenever we encounter a star, we count it and then remove it from the picture.

For instance, black out all the `-` cells belonging to the star we wish to remove by replacing them with `#`.

Counting Stars — Example

Let's start off with this grid.

We iterate through row-wise followed by column-wise to find the first - cell.

```
#####  
## - - ## - ##  
# - - #####  
# - - - # - - ##  
#### - - - ####  
#####  
#### - - ## - #  
##### - #####  
#####
```

Counting Stars — Example

Encountered first star.

Black out all its -

```
#####  
##- -##-##  
#- -#####  
#- - -#- -##  
###- - -###  
#####  
###- -##-#  
####-####  
#####
```



```
#####  
######-##  
#####  
#####  
#####  
#####  
###- -##-#  
####-####  
#####
```


Counting Stars — Example

Encountered **second**
star.

Black out all its -

```
#####  
######-##  
#####  
#####  
#####  
###########  
#####  
###- -##-#  
####-#####  
#####
```



```
#####  
##########  
#####  
#####  
#####  
###########  
#####  
###- -##-#  
####-#####  
#####
```

Counting Stars — Example

Encountered **third** star.

Black out all its -

```
#####  
########  
#########  
#####  
#####  
#####  
#####  
###- - ## - #  
##### - #####  
#####
```



```
#####  
########  
#####  
#####  
#####  
#####  
#####  
####### - #  
#####  
#####
```

Counting Stars — Example

Encountered **fourth** star.

Black out all its -

```
#####  
########  
#########  
##########  
########  
#####  
########-##  
#####  
#####
```



```
#####  
########  
#########  
##########  
########  
#####  
##########  
#####  
#####
```

Counting Stars — Example

In total, there are 4 stars.

As a graph problem:

- Task: count the number of CC
- Cells \rightarrow vertices
- Edges \rightarrow between adjacent cells that are -
- Flood filling \rightarrow DFS

```
#####  
##########  
#####  
#####  
#####  
#####  
#####  
#####  
#####  
#####
```

Example Problem 2: [UVa 469](https://uva.onlinejudge.org/external/4/469.pdf)

URL: <https://uva.onlinejudge.org/external/4/469.pdf>

You are given a N by M grid of cells.

Each cell is either **L** or **W** representing *land* and *water* respectively.

Answer K queries of the following:

- Given a **W** cell, how many **W** cells are connected to it.

UVa 469—Example

Diagonal cells are considered adjacent

Red: 12

Purple: 1

Blue: 4

Orange : 1

```

LLLLLLLLLL
LLWWLLWLL
LWWLLLLLL
LWWWLWWLL
LLLWWWLLL
LLLLLLLLLL
LLLWLLWL
LLWLWLLL
LLLLLLLLLL
  
```

UVa 469—Approach

If we group all connected **W** cells as 1 water body.

We need to obtain the size of the water body each **W** cell is in.

```
LLLLLLLLLL
LLWWLLWLL
LWWLLLLLL
LWWWLWWLL
LLLWWWLLL
LLLLLLLLLL
LLLWWLWL
LLWLWLLLL
LLLLLLLLLL
```

UVa 469—Graph Modelling

- Vertex: Cells of the grid
- Edge: 8 radial directions
 - Draw an edge between 2 adjacent **W** cells
- Number of vertices in the connected component: Size of the water body

```
LLLLLLLLLL
LLWWLLWLL
LWWLLLLLL
LWWWLWWLL
LLLWWWLLL
LLLLLLLLLL
LLLWLLWL
LLWLWLLL
LLLLLLLLLL
```


UVa 469—Implementation

What algorithm can we use to solve this?

- a) Depth First Search
- b) Breadth First Search
- c) Either one

Deeper Stuffs about Topological Sort

Recap: Topological sort

Topological sorted order on a **DAG** is

- A **linear ordering** of all the DAG's vertices
- Criteria:
 - Graph must be a DAG
 - For every edge $u \rightarrow v$ in the graph, u must come before v in the topological sort!

Test yourself

Given the following toposort

1, 2, 3, 4, 5

Is it guaranteed that vertex 1 has an outbound edge to 2?

Test yourself

Given the following toposort

1, 2, 3, 4, 5

Is it guaranteed that vertex 1 has an outbound edge to 2?

Answer: No! The rule just say that if $u \rightarrow v$, then v must come after u in the toposort. However, **the converse is not true!** Knowing that v come after u in the toposort is insufficient for us to conclude that $u \rightarrow v$.

So in this case, all we can say is 1's neighbours are on its right in the toposort but we can't tell who they are if they exist.

Recap: Toposort implementation — DFS subroutine

Only 1 new line added to DFS!

```
void toposort(int v_id) {  
    if (visited[v_id]) return;  
    visited[v_id] = true;  
    for (auto &nb_id : AL[v_id]) {  
        toposort(nb_id);  
    }  
    tps.push_back(v_id);  
}
```

} Standard DFS code

// Append to toposorted list

Recap: Toposort implementation — Main routine

Our main procedure looks like this!

```
...  
/* Iterate over each vertex v_id in graph */  
for (int v_id = 0; v_id < AL.size(); ++v_id) {  
    toposort(v_id); // Call toposort on vertex v_id  
}  
/* Reverse toposorted list tps */  
reverse(tps.begin(), tps.end());  
...
```

Assume vertices in graph numbered from 0 to V

Test yourself!


Why do we need to call toposort on every vertex?

Test yourself!


Why do we need to call toposort on every vertex?

Answer: Because a given vertex might not be able to reach every other vertex in the DAG! In the generalised case, want the toposort on the entire graph, not just on a single vertex.

E.g. Consider the graph on the right.




Recap: Topological sort V.S. DFS



DFS visitation order: 0, 2, 3, 4, 7, 1, 5, 6

Is this what we want?

Recap: Topological sort V.S. DFS




DFS visitation order: 0, 2, **3**, 4, 7, **1**, 5, 6

Is this what we want?

Answer: No! Because 1 should come before 3!
This proves that DFS visitation order is not the same as toposort order!


Recap: Topological sort V.S. DFS



DFS visitation order: 0, 2, 3, 4, 7, 1, 5, 6

Toposort order: 5, 6, 1, 0, 2, 4, 7, 3


Recap: Topological sort properties



Toposort order:

5, 6, 1, 0, 2, 4, 7, 3

With this sequence, we can also redraw the graph in a linear fashion!



Recap: Topological sort properties

Corollary

Realise this also means that in the DAG, if vertex u **has a path** to vertex v , then v **must also be on the right** of u in the toposort list!

Test yourself!

Can undirected graphs have a topo sort?


Hint: Try running it on an undirected graph and a directed cyclic graph

Question 3

- Please review <https://visualgo.net/en/dfsbfbs> (either DFS or BFS version)
- The modified DFS or modified BFS (Kahn's) topological sort algorithm only gives one valid topological ordering
- How can we find **all possible valid topological orderings** for a given DAG?

Valid topological orderings


Recall: Toposort of a graph is not unique. Multiple possible sorted orders exist! However, a DAG must have **at least** one topological ordering.



What are all of the toposorts for this graph?

Valid topological orderings

Recall: Toposort of a graph is not unique. Multiple possible sorted orders exist! However, a DAG must have **at least** one topological ordering.




What are all of the toposorts for this graph?

Answer:

Realize that:


- 1, 0, 2, 3
- 0, 1, 2, 3
- 0, 2, 1, 3
- Every other vertex point to 3 and therefore 3 **must be** the last in the toposort
- It doesn't matter how 1 is ordered relative to 0,2,3 in the toposort. Why?

Valid topological orderings



How many valid topological orderings does this graph have?

Valid topological orderings



How many valid topological orderings does this graph have?

Answer: 1008. Don't try manually counting!

How did we derive this? We ran the code on the next slide :O

Valid topological orderings

```
#include <bits/stdc++.h>
using namespace std;
typedef pair<int, int> ii;
int main() {
    int V = 9; // 9 vertices
    vector<ii> EL; // Store the graph as edge list
    EL.emplace_back(0,1); EL.emplace_back(0,5); EL.emplace_back(1,4); EL.emplace_back(2,0); EL.emplace_back(2,1);
    EL.emplace_back(2,4); EL.emplace_back(5,3); EL.emplace_back(6,7); EL.emplace_back(6,8);
    int p[10];
    for (int i = 0; i < V; ++i) p[i] = i;
    int ans = 0;
    do {
        bool valid = true;
        for (int i = 0; i < (int)EL.size(); ++i) {
            auto [u, v] = EL[i];
            int pos_u, pos_v;
            for (pos_u = 0; pos_u < V; ++pos_u) if (p[pos_u] == u) break;
            for (pos_v = 0; pos_v < V; ++pos_v) if (p[pos_v] == v) break;
            if (pos_u > pos_v) valid = false;
        }
        if (valid) ans++;
    } while (next_permutation(p, p+V));
    printf("There are %d valid topological orderings\n", ans);
    return 0;
}
```

More on how this works in a bit!

Valid topological orderings—Starting point

Question to ask as a starting point:

What kind of DAG has the

- Minimum
- Maximum

number of valid topological orderings?

Valid topological orderings—Starting point

What's the **minimum** number of valid topological ordering for a DAG?

Valid topological orderings—Starting point

What's the **minimum** number of valid topological ordering for a DAG?

Answer: 1. Because a DAG must have at least 1 valid topological ordering.

Valid topological orderings—Starting point

What kind of DAG has exactly 1 topological ordering?

Valid topological orderings—Starting point

What kind of DAG has exactly 1 topological ordering?

Answer: A linked list!

Valid topological orderings—Starting point

What's the **maximum** number of valid topological orderings for a DAG?

Valid topological orderings—Starting point

What's the **maximum** number of valid topological orderings for a DAG?

Answer: A topological ordering is a *permutation* of V vertex numbers. So maximum number of toposorted orders is $V!$

Valid topological orderings—Starting point

How can we construct a graph such that every permutation is a valid toposort order?

Valid topological orderings—Starting point

How can we construct a graph such that every permutation is a valid toposort order?

Answer: A disconnected graph! i.e. a graph with V edgeless vertices. Any relative ordering of vertices is a valid toposort since there are no edges.

Valid topological orderings—Approach

Back to the question: How can we find all possible valid topological orderings for a given DAG?

Worse case: $V!$ topological orderings.

Brute Force Checking: We can iterate through all $V!$ permutations and check each one if they are valid.

Valid topological orderings—Validating permutation

How to check if a permutation is valid?

For every edge in the graph from $u \rightarrow v$, we make sure that u appears before v in the permutation.

Invalid otherwise.

Valid topological orderings—Naive approach

For **each edge** $u \rightarrow v$ in the graph, how to check if u comes before v in the permutation?

Naive approach: Iterate the entire permutation and check if we encounter u before we encounter v .

- $O(V)$ time **per edge**
- $O(VE)$ time **over entire graph**

Valid topological orderings—Naive approach

Running this naive approach across all $V!$ possible permutations therefore incurs a total time complexity of

$$O(V! \times (VE))$$

Recall the code we presented a few slides ago to calculate the 1008 possible toposorts of the example graph?

It uses exactly this naive approach!

Valid topological orderings—Slightly better approach

A slightly better approach:

For each permutation, we first pay a $O(V)$ penalty to preprocess it by using a table to map each vertex to their position in the permutation.

Thereafter, for **each edge** $u \rightarrow v$, we just need to check in the permutation if the position of u comes before the position of v .

- $O(1)$ time **per edge**
- $O(V + E)$ time **over entire graph** (including preprocessing time)

Valid topological orderings — Slightly better approach

Example:

Say for instance we have the following permutation to validate:

1, 5, 6, 2, 7, 4, 3

So we first build the table on the right.

Then to validate $5 \rightarrow 7$ for instance, we check if 5 comes before 7 in the permutation by verifying $1 < 4$.

Vertex	Permutation position
1	0
2	3
3	6
4	5
5	1
6	2
7	4

Valid topological orderings—Slightly better approach

Running this slightly better approach across all $V!$ possible permutations therefore incurs a total time complexity of

$$O(V! \times (V + E))$$

Valid topological orderings—Summary

Time complexities of various solutions to the problem:

- Naive approach (demo code): $O(V! \times VE)$
- Slightly better approach: $O(V! \times (V + E))$
- Backtracking: $O(V! \times V)$
- Backtracking with bitmask: $O(V!)$ (Beyond CS2040C)

Question 4

The modified BFS (Kahn's) topological sort algorithm is actually quite interesting! Read more about the details on [Wikipedia](#).

On VisuAlgo, we used a `std::queue` for the underlying data structure in the modified BFS. What if we replaced it with

- A stack `std::stack`
- A priority queue `std::priority_queue` or BBST `std::set` (hopefully you now know that this is also a valid ADT Priority Queue implementation)
- A Hash Table `std::unordered_set`

Toposort implementation — Kahn's algorithm

Pseudocode (modified from VisuAlgo)

```
void kahns_algo() {  
    Q = queue(); // Initialize empty queue  
    for (each vertex v in graph) { // Pre-populate Q  
        if (v has no incoming edge) Q.push(v);  
    }  
    while (!Q.empty()) {  
        u = Q.front(); Q.pop(); // Dequeue from Q  
        tps.push_back(u);      // Append to toposort list  
        for (each neighbor v of u) {  
            remove_edge(u,v); // Remove edge u→v from graph  
            if (v has no incoming edge) Q.push(v);  
        }  
    }  
}
```


Toposort implementation — Kahn's algorithm

Although trickier to implement than modified DFS, Kahn's algorithm is actually much more intuitive and its steps can be easily understood as follows.

So long as there are vertices in the graph:

- Greedily select the vertices with no *incoming* edges
- Add them to toposort list
- Remove these vertices and their edges from the graph

Kahn's algorithm: Implement using Stack

If we replace the `std::queue` with a `std::stack`

- We will also get another valid topological ordering
- Topological ordering in LIFO instead of FIFO order in case of ties (more than one vertices with 0 in-degree)

Kahn's algorithm: Implement using Priority Queue

If we replace the `std::queue` with a (max/min) `std::priority_queue` or `std::set`:

- We will also get another valid topological ordering
- Topological ordering in (max/min) priority order in case of ties (more than one vertices with 0 in-degree)

There are a few topological sort problems that require this modification!

Kahn's algorithm: Implement using Hash Table

If we replace the `std::queue` with a `std::unordered_set`:

- We will also get another valid topological ordering
- Topological ordering in unspecified order in case of ties (more than one vertices with 0 in-degree)

Wow Kahn's algorithm works even when the vertices to be removed are not ordered in any way! :O

Kahn's algorithm: Observations

- Basically Kahn's algorithm is quite versatile as it just need the data structure to be a set S for containing vertices
- Realize that at any one point in time, S always contains vertices with 0 *in-degree* (i.e. no incoming edges). That's actually its sole purpose
- We saw that S can be FIFO ordered, LIFO ordered, fully ordered, or not ordered at all! The order doesn't seem to matter!

Kahn's algorithm: Observations

Why doesn't it matter how S orders its vertices?

- The order of vertices in S just captures the sequence in which the 0 in-degree vertices are to be removed and appended to toposort list
- Realize that 0 in-degree vertices do not point to each other and therefore their *relative ordering* in the toposort **does not matter!**

Kahn's algorithm: Jin's informal proof of correctness

- [1] Given: A DAG must have 0 in-degree vertices, otherwise it will not be a DAG because it have cycles.
- [2] Given: A DAG is still a DAG after having vertices removed because cycles cannot be introduced from edge removals.
- [3] We have a graph G which is a DAG.
- [4] After removing all 0 in-degree vertices from a G , we have a new graph G^* .
- [5] From [2], G^* must still be a DAG.
- [6] From [1], G^* must have its own set of 0 in-degree vertices.
- [7] If we repeatedly remove 0 in-degree vertices, we will never encounter a vertex which had an outgoing edge to a vertex we removed previously. i.e. We will never encounter violation.
- [8] From [7], therefore topological order is guaranteed via the sequence of 0 in-degree vertex removals. *Q.E.D*

Toposort Sort—Protip

When coming up with a toposort by hand, it's a lot simpler to use Kahn's algorithm. You may want to do this for VisuAlgo quizzes.

To carry out Kahn's algorithm, just follow these steps:

1. In any order, find vertices with no *incoming edges*
2. As soon as you find one
 - a. Append it to toposort list
 - b. Remove it and all its edges
3. Repeat steps 1 & 2 until no more vertices are left in the graph

Facebook Privacy Setting

CS2010 Finals AY2013/2014 Sem 1

$O(V+E)$ solution

$O(k)$ solution

Problem statement

- You have a graph of V vertices, E edges
- You are given a pair of vertices i and j .
 - Compute whether vertex i and j are *at most 2* edges apart.
i.e. degree of separation is 2
- Given: The friend list of profile i is stored in `adjList[i]` and sorted ascending
- $O(V+E)$ for 7 marks
- $O(k)$ for 19 marks
 - k is sum of number of adjacent vertices of i and j

How to get $O(V+E)$ solution?

Bellman Ford?

- $O(VE)$

Dijkstra?

- $O((V+E)\log V)$



JAKE-CLARK.TUMBLR

$O(V+E)$ solution

Breadth First Search!

- Start from vertex i , compute shortest path from i to every other vertex
- Check if $\text{distance}(i, j) \leq 2$

Can we do better?

Observation

We only need $\text{distance}(i, j) \leq 2$

3 cases:

1. $\text{distance}(i, j) = 0$ $i = j$
2. $\text{distance}(i, j) = 1$ j is a neighbour of i
3. $\text{distance}(i, j) = 2$ j is a neighbour of a neighbour of i

Case 1: $\text{distance}(i, j) = 0$


This is trivial, just $O(1)$ to check if $i = j$

Case 2: $\text{distance}(i, j) = 1$

To check if j is a neighbour of i , we search for j in each of i 's k_i neighbours.

Since given that friends list is sorted, binary search will take $O(\log k_i)$

Is j within green?




Green: distance = 1

Case 3: $\text{distance}(i, j) = 2$

To check if j is a neighbour of a neighbour of i , we potentially traverse the entire graph and so this will take $O(V + E)$.

Can we do better than this?

Is j within orange?



Green: distance = 1
Orange: distance = 2

A relevant question


You are in **NUS** now. (West)

Your partner just touched down at **Changi Airport**. (East)

What should y'all do in order to meet each other in the shortest possible time?

A relevant question

- A. Meet at NUS
- B. Meet at Changi
- C. Meet somewhere in the middle
- D. Use video call
- E. This is a scam. I don't have a partner




A relevant question

- A. Meet at NUS
- B. Meet at Changi
- C. Meet somewhere in the middle
- D. Use video call
- E. This is a scam. I don't have a partner

Case 3: $\text{distance}(i, j) = 2$


Is your partner in orange?



Green: distance = 1
Orange: distance = 2

Case 3: $\text{distance}(i, j) = 2$


If vertex i and j are distance 2 away from each other, what do you realize?



Case 3: $\text{distance}(i, j) = 2$

If vertex i and j are distance 2 away from each other, what do you realize?

Answer: There must be at least one common friend!



Transformed Problem

Given 2 integer arrays with sizes up to N each, find whether there exist an integer that is in both arrays.

Approaches:

1. For every number in one array, search for a correspondence in the other array $O(N^2)$
2. Sort both arrays $O(N \log N)$, then iterate through both arrays with 2 pointers to look for a corresponding pair of numbers $O(N)$

Facebook Privacy Setting

Realize that the transformed problem is essentially what we want to solve in the original problem!

Solution 2 (previous slide) would just take $O(k)$ in the original problem since we are already given that the friends list is sorted in the adjacency list! We just need to iterate down the friends list of i and j once!

Moral of the Story

Don't judge a book by its cover!

What might appear like a Graph question, might not actually be a Graph question.

What might appear not a Graph question, might actually be a Graph question.

Questions?